


Columnar Storage and List-based Processing for Graph Database Management System

Pranjal Gupta, Amine Mhedhbi, Semih Salihoglu

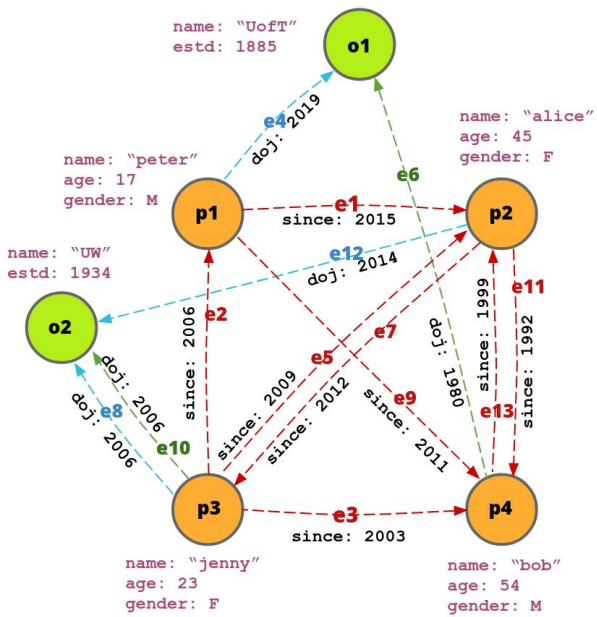


UNIVERSITY OF
WATERLOO



Graph Database Management Systems

Read-optimized analytical systems that can perform fast many-to-many joins.



Property Graph Data Model

p1	→	e1,p2	e9,p4	e4,o1		
p2	→	e7,p3	e11,p4			
p3	→	e2,p1	e3,p4	e5,p2	e8,o2	e10,o2
p4	→	e13,p2				

Forward Adjacency Lists

p2	name	"alice"	age	45	gender	F
p4	name	"bob"	age	54	gender	M
...						
o2	name	"UW"	estd	1934		

Row-oriented Property Store

```
MATCH
(a:PERSON) - [e:FOLLOWS] → (b:PERSON)
WHERE
a.age > 30 & e.since > 2000
RETURN ...
```

Cypher Query

Fast *many-many joins* are attributed to the adjacency lists data structure that indexes neighbours of all vertices in both directions, by default.

Analogous to *Indexed Nested Joins in RDBMS*.

Relevance of Columnar Techniques for GDBMSs

Columnar RDBMSs are optimized for read heavy workloads that requires reading large volumes of data and processing it.

Storage

p4	`bob`	54	TO
p2	`jenny`	23	TO
p3	`alice`	45	VC
p1	`peter`	17	VC
	PERSON .name	PERSON .age	PERSON .city

Compression

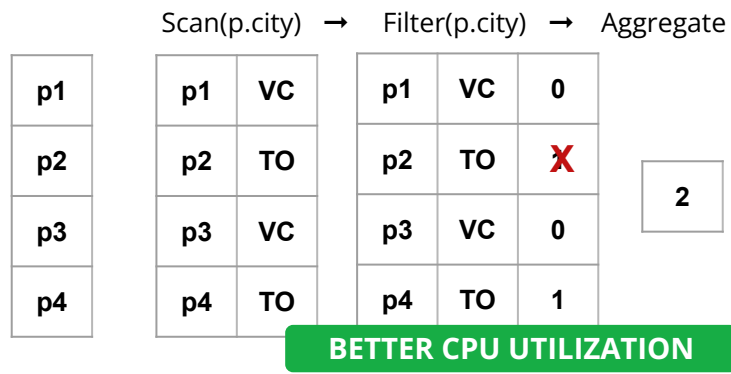
0011b

VC=0, TO=1

Dictionary-encoded PERSON.city

Block-based processing

```
SELECT count(*)
FROM PERSON p
WHERE p.city = `TO`
```



GDBMSs can benefit from these techniques.

However ! Not directly applicable because of major differences in the nature of the graph data and access patterns of GDBMSs


Difference in Access Patterns & Workloads (1)

Columnar RDBMSs

Queries are predominantly operations over large chunks *sequential* column data.

Eg. Group by and Aggregates, Filter etc.

```
SELECT p.city, count(*)
FROM Person p
GROUP BY p.city
```





GDBMSs

Queries requires *random* access to large chunks of data.

Eg. n-hop queries

```
MATCH (a:Person)-[e:Knows]->(b:Person)
WHERE a.name = 'alice'
RETURN b.age
```




Cannot use vanilla columnar compression techniques directly.
Because these techniques require decompressing an entire block sequentially

Difference in Access Patterns & Workloads (2)


Columnar RDBMSs

Primary Key - Foreign Key joins are prevalent.
Intermediate join results do not grow.



GDBMSs

Many-to-many joins are prevalent.
Intermediate results grow with **value replication**.



alice	k1	k11
alice	k1	k12
alice	k1	k13
...
alice	k3	k32
alice	k3	k33

Intermediate tuples

Vanilla Block-based Processing do not avoid value replication and hence, can be inefficient.

Research Contributions

Columnar RDBMSs		GDBMSs
<div>STORAGE</div> <div>Vanilla append-only Columns</div>	<div>DIRECT ADAPTATION</div>	<div>Vanilla columns for vertex properties.</div> <div>NOVEL</div> <div>Single-index edge property pages.</div>
<div>COMPRESSION</div> <div>NULL Compression using bit-strings <i>[Abadi, CIDR 2007]</i></div>	<div>DIRECT ADAPTATION</div> <div>NOVEL</div> <div>NOVEL</div>	<div>Dictionary Encoding</div> <div>NULL Compressing using Jacobson's bit vector index</div> <div>Compression using new edge ID scheme</div>
<div>QUERY PROCESSING</div> <div>Block-based processing</div>	<div>NOVEL</div>	<div>List-based processing</div>


Single-indexed Edge Property Pages

Observation

Edge and Vertex properties are read in the same order as edges in the adjacency lists.

```
MATCH (a:PERSON) - [e:FOLLOWS] → (b:PERSON)
WHERE a = p3
RETURN e.since, b.age
```

[Scan a] → [Filter a = p3] → [Join a with b] → RETURN e.since, b.age



REQUIRED PROPERTY

Edge properties should be stored & read sequentially in the order edges are stored in the adjacency lists.

However ! We cannot do the same for vertex properties.

e1	2015		
e11	1992		
e3	2003		
e7	2012		
e5	2009		
e9	2011		
e2	2006		
e13	1999		
KNOWS since			

	p4	54
	p2	23
	p3	45
	p1	17
	PERSON age	

Columnar Storage

Simple append-only Columns:

Can be used to store vertex properties.

Direct access to the property value using offset in column.

Suboptimal with edge properties, as does not satisfy our required of sequential reads.

NON-SEQ READS

p4	54
p2	23
p3	45
p1	17

PERSON
age

Doubly-indexed Edge Property Lists:

Inspired by Adjacency Lists.

SEQ READS

2x DATA DUPLICATION

p1 →	2015	2011	
p2 →	2012	1992	
p3 →	2006	2003	2009
p4 →	1999		

KNOWS
since

FWD

p1 →	2006		
p2 →	2015	2009	1999
p3 →	2012		
p4 →	2003	2011	1992

BWD

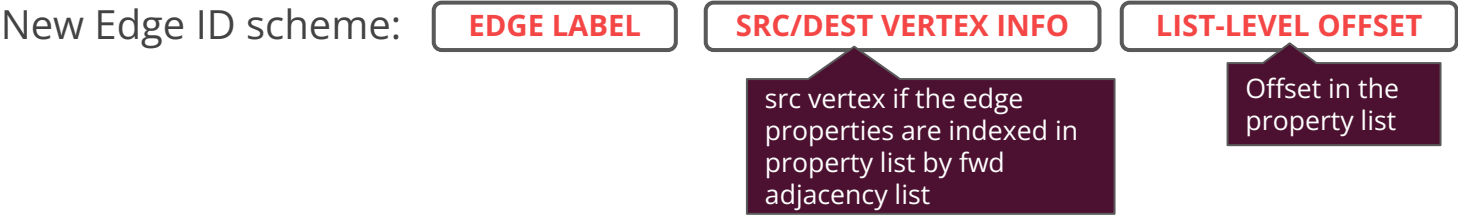

Single-indexed Edge Property Lists.

Keep property lists in only one direction.

Provide for random access of the property when edges are read from adjacency list of opposite direction.

SEQ READS IN ONE DIRECTION

For an edge, requires an offset in the appropriate property list to access its property:



Optimized for updates: **Single-indexed Edge Property Pages**

Evaluation

- Datasets: LDBC100, FLICKR, WIKI
- Queries: n-Hop with either Filter or Aggregation on the last Join
- Configurations: **COL_E** - Simple append-only columns for storing edge properties
PAGE_E - Single-directional edge property pages for edge properties

		LDBC100		WIKI		FLICKR	
		1H	2H	1H	2H	1H	2H
P_F	COL _E	0.55	65.22	2.97	42.92	1.88	888.30
	PAGE _P	0.16 3.4x	34.22 1.9x	0.96 3.1x	16.48 2.6x	0.42 4.5x	189.39 4.7x
P_B	COL _E	1.23	131.01	6.33	99.28	2.40	1009.84
	PAGE _P	1.29 0.9x	134.43 1.0x	6.10 1.0x	91.75 1.1x	2.25 1.1x	1183.14 0.9x



NULL Compression using Jacobsons' bit vector index

Observation

Properties on edges and vertices can be very sparse.

The degree of most vertices are super small. (by power law)

Reading vertex property (in FILTER) and adjacency list (in JOIN) is like reading small chunk of data followed by random access.



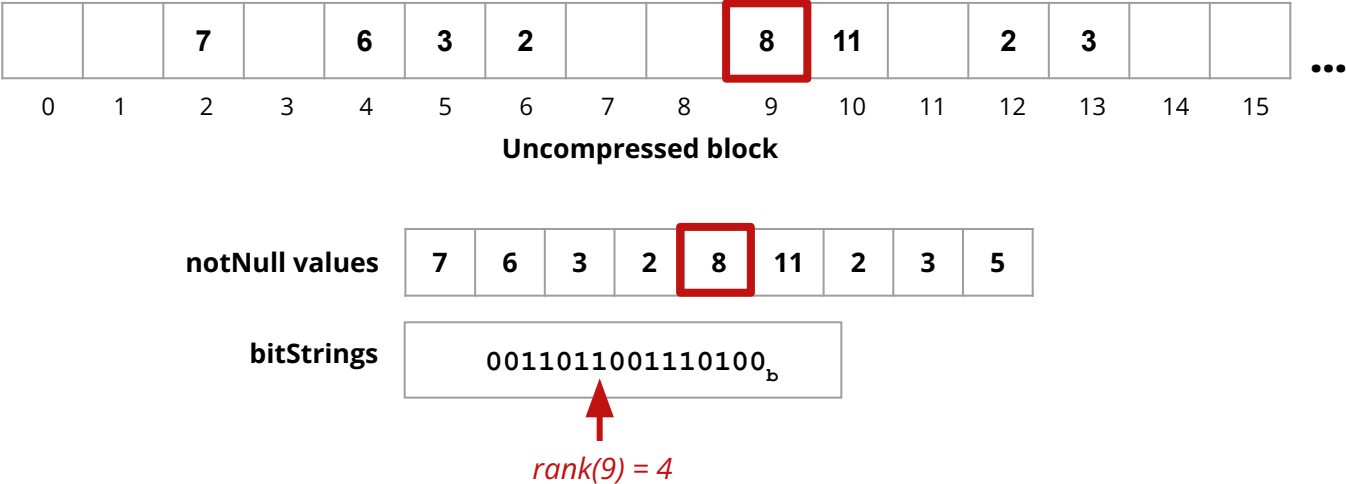
REQUIRED PROPERTY

If compression is used, decompressing arbitrary data elements in a compressed block should happen in constant time.

NULL Compression

Existing solution to compress columns using bit-strings. [Abadi, CIDR 2007]



Suitable only for sequential access of column. Do not support reading from arbitrary offsets of columns.



Using Jacobson's bit vector index

Solution: Calculate the rank of an offset in constant time.

Use Jacobson's rank index. [Jacobson, FOCS 1989]



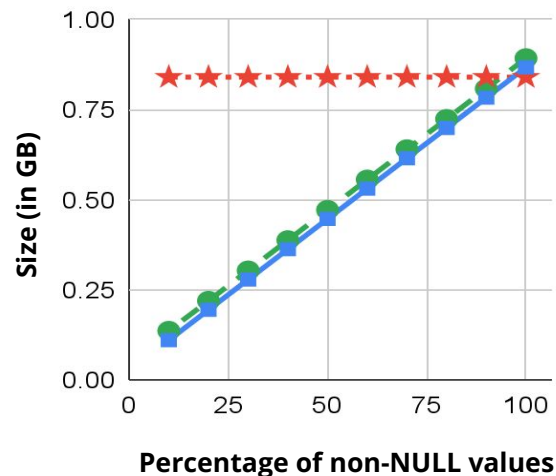

chunkIdx = 2

position of element in chunk (i) = 1

position in notNull values array = 4 + 0 = 4

Evaluation

- Datasets: 1 vertex property column with 220M entries) with different %s of NULL values
- Queries: 1-Hop with Property access
- Configurations: ★ Uncompressed ● J-NULL ■ Vanilla-NULL




List-based Processing

Volcano-styled Processing

```
MATCH (a:PERSON) - [e:FOLLOWS] → (b:PERSON)
WHERE a = p3
RETURN e.since, b.age
```

[Scan a] → [Filter a = p3] → [Join a with b] → RETURN e.since, b.age



p3	e3	p2	23	2006
a	e	b	b.age	e.since

Bad Cache Locality !

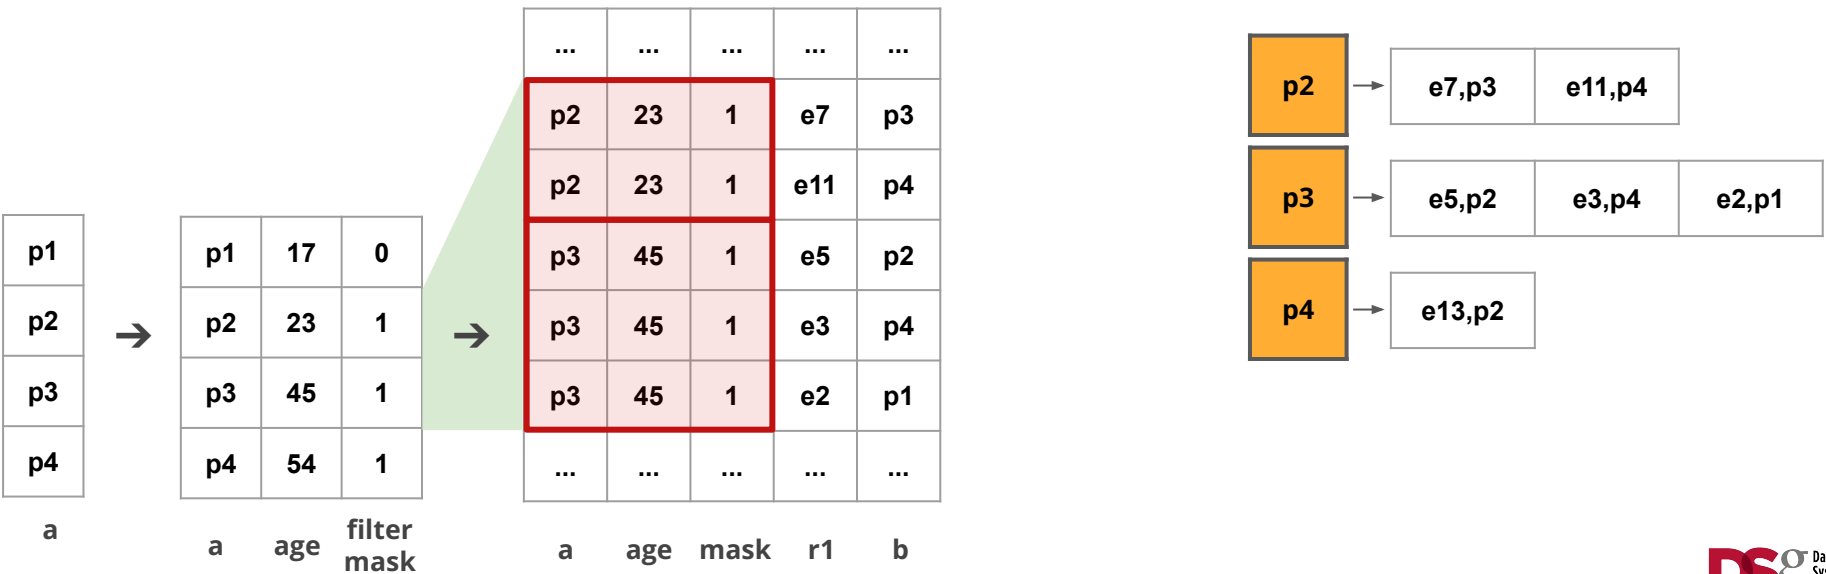
Do not harness the fact that adjacency lists and edge properties are stored sequentially in memory.

e1	2015		
e11	1992		
e3	2003		
e7	2012		
e5	2009	p4	54
e9	2011	p2	23
e2	2006	p3	45
e13	1999	p1	17
KNOWS since		PERSON age	

Shortcomings of Block-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] -> (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] -> (c:PERSON)
WHERE a.age > 20
RETURN ...
```


[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN ...



Shortcomings of Block-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] -> (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] -> (c:PERSON)
WHERE a.age > 20
RETURN ...
```

[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN ...



Shortcomings of Block-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] → (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] → (c:PERSON)
WHERE a.age > 20
RETURN ...
```

[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN ...

Shortcoming #1

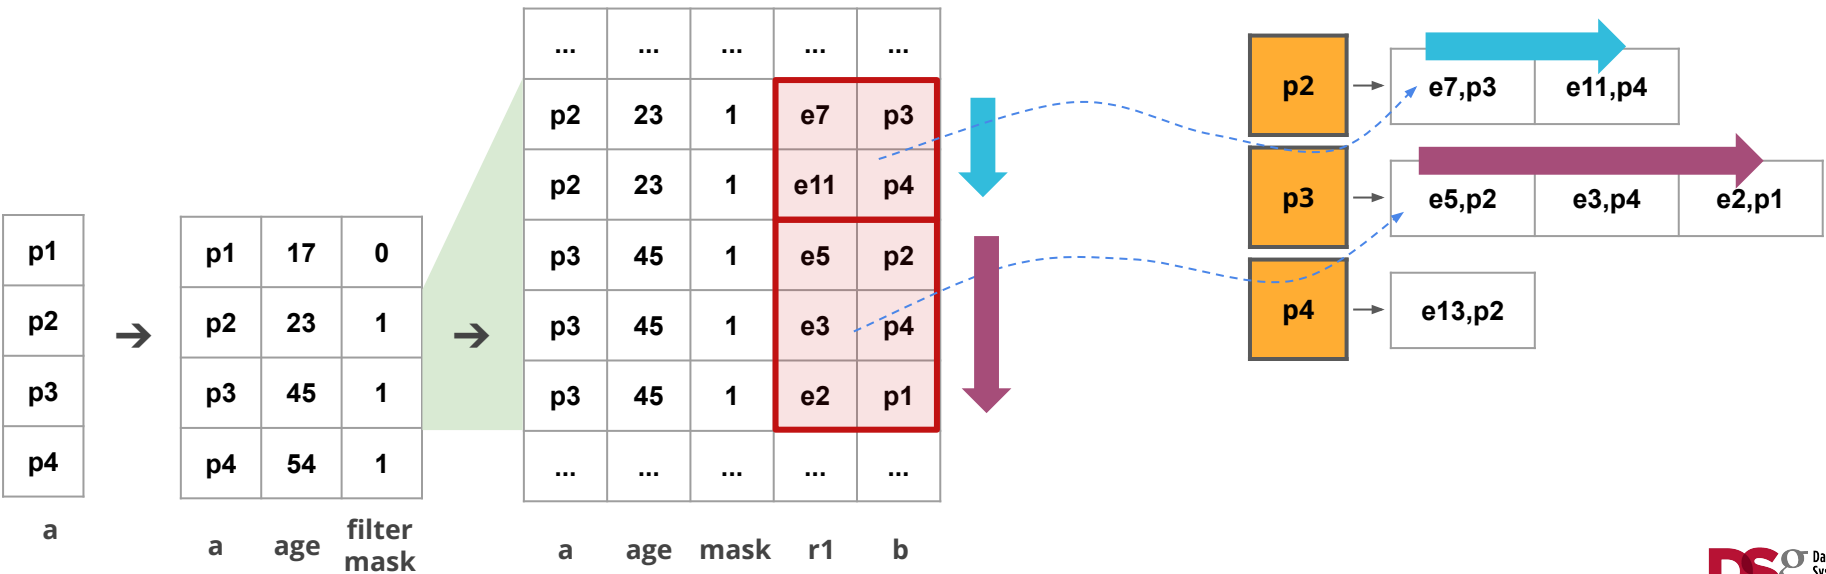
High amount of data replication in intermediate data chunks.

Because intermediate data chunk represents a set of **flat values** using 1 group of vectors.

Shortcomings of Block-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] -> (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] -> (c:PERSON)
WHERE a.age > 20
RETURN ...
```

[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN ...



Shortcomings of Block-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] → (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] → (c:PERSON)
WHERE a.age > 20
RETURN ...
```

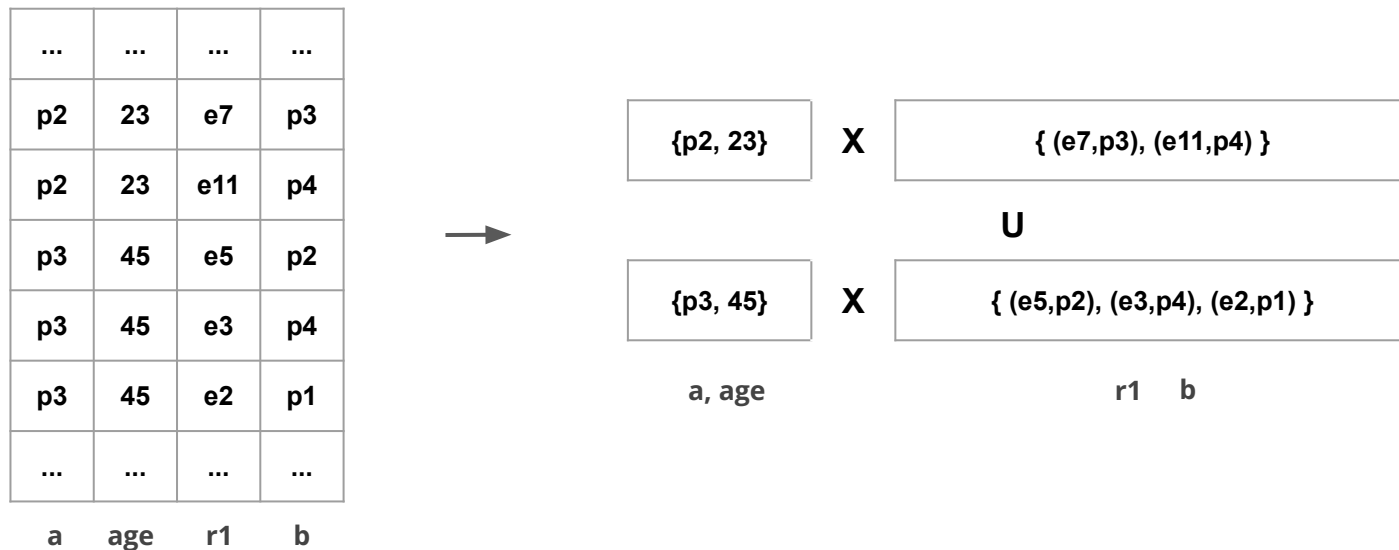
[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN ...

Shortcoming #2

Copies exact replicas of adjacency lists to intermediate data chunk.

List-based Processing

Factorized representation of intermediate data [Olteanu, SIGMOD Rec. 2016].



List Groups


- Instead of 1 group of vectors, **multiple groups** that either be a single tuple or list of tuples.
- Instead of fixed-size vectors, **variable-size vectors** that depend on adjacency list sizes.
- Avoid materializing adjacency lists in list groups

List-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] → (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] → (c:PERSON)
WHERE a.age > 20
RETURN ...
```

[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN

....



List-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] → (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] → (c:PERSON)
WHERE a.age > 20
RETURN ...
```

[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN
....

LG1

p1	17	0
p2	23	1
p3	45	1
p4	54	1


a age mask

curridx = -1

List-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] -> (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] -> (c:PERSON)
WHERE a.age > 20
RETURN ...
```

[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN
....




Performance speedups upto 19x
on LDBC100

List-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] -> (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] -> (c:PERSON)
WHERE a.age > 20
RETURN ...
```

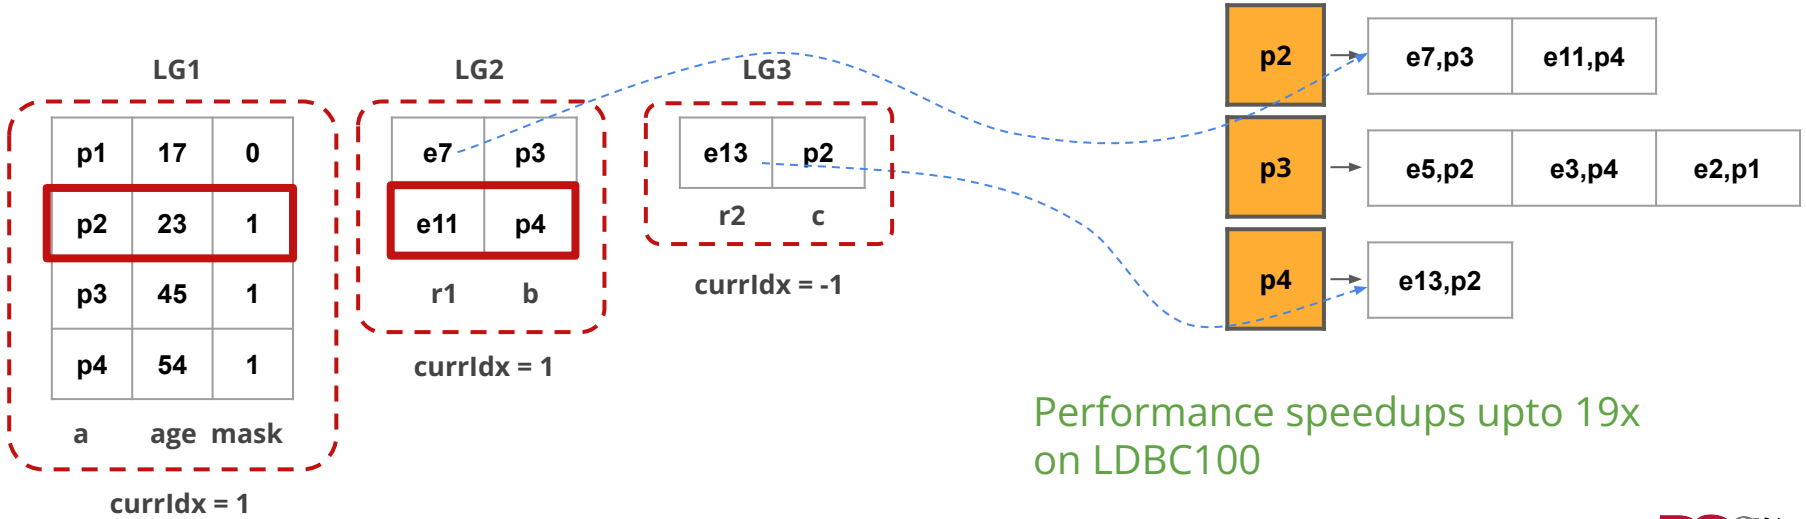
[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN
....



List-based Processing

```
MATCH (a:PERSON) - [r1:FOLLOWS] -> (b:PERSON)
      (b:PERSON) - [r2:FOLLOWS] -> (c:PERSON)
WHERE a.age > 20
RETURN ...
```

[Scan a] → [Filter a.age > 20] → [Join a with b] → [Join b with c] → RETURN
....



Evaluation

- Queries: n-Hop with either Filter or Aggregation on the last Join
- Configurations: **GF-CV** - Column-oriented, Volcano-styled processing
GF-CL - Column-oriented, List-based processing
- Datasets: LDBC100, FLICKR, WIKI

			1-hop	2-hop	3-hop
LDBC100	FILTER	GF-CV	24.6	1470.5	40252.4
		GF-CL	7.7	116.2	2647.3
	COUNT(*)	GF-CV	13.4	241.9	6947.3
		GF-CL	4.2	18.9	357.9
			3.2x	12.7x	15.2x
			3.2x	12.8x	19.4x

References

Column Stores for Wide and Sparse Data.

Daniel J. Abadi. CIDR 2007

Column-Stores vs. Row-Stores: How Different Are They Really?

Daniel J. Abadi, Samuel Madden, and Nabil Hachem. SIGMOD 2008

Factorized Databases.

Dan Olteanu and Maximilian Schleich. SIGMOD Record 2016.

C-store: A Column-Oriented DBMS.

Mike Stonebraker, et. al.. Making Databases Work 2019

MonetDB: Two Decades of Research in Column-oriented Database Architectures.

Stratos Idreos et. al.. IEEE Data Engineering Bulletin 2012

Space-efficient Static Trees and Graphs.

Guy Jacobson. FOCS 1989


Vectorwise: Beyond Column Stores.

Marcin Zukowski and Peter A. Boncz. IEEE Data Engineering Bulletin 2012

Volcano - An Extensible and Parallel Query Evaluation System

Goetz Graefe. IEEE Transactions on Knowledge and Data Engineering, TKDE 1994

Visit us at: <http://graphflow.io>



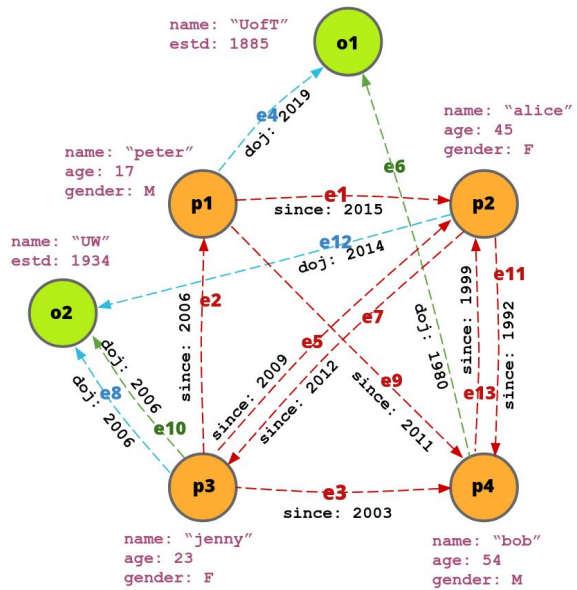
End.

Questions?

Extras

Graph Database Management System (GDBMS)

GDBMSs are systems are *read-optimized* systems that are known to support read heavy workloads that predominantly contains large number of *many-to-many joins*.

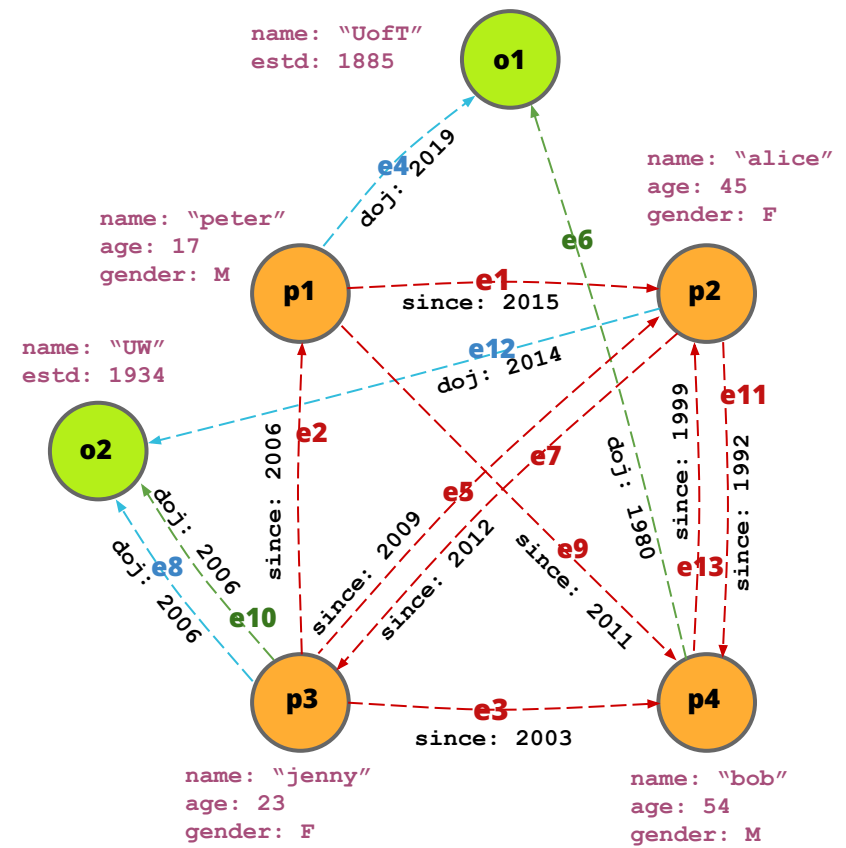


p1 →	e1,p2	e9,p4	e4,o1		
p2 →	e7,p3	e11,p4			
p3 →	e2,p1	e3,p4	e5,p2	e8,o2	e10,o2
p4 →	e13,p2				

Adjacency Lists

Applications like fraud detection, recommendation systems, social networks.

Graph Databases 101: Model



- Vertices
- Vertices have a **vertex type**

ORG


PERSON
- Directed **edges** connect vertices
- Edges have a **label**

FOLLOWS

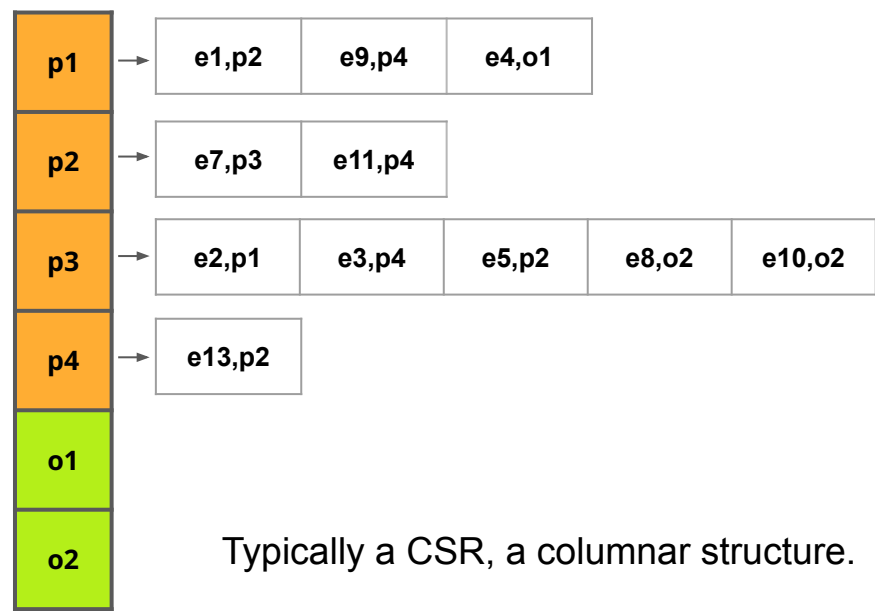
WORKAT

STUDYAT
- Schemaless, but Arbitrary key-value **properties** on edges and vertices

Graph Databases 101: Storage

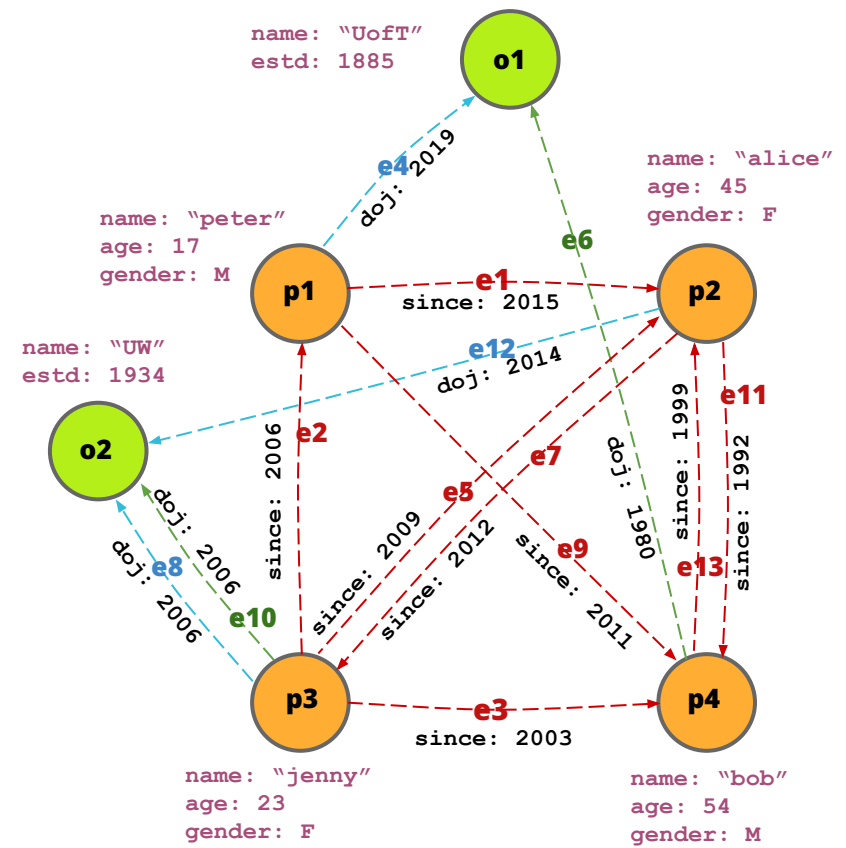


Storing the graph topology:



Typically a CSR, a columnar structure.

Graph Databases 101: Storage



Storing the properties:

p2

name	"alice"	age	45	gender	F
------	---------	-----	----	--------	---

p4

name	"bob"	age	54	gender	M
------	-------	-----	----	--------	---

...

o2

name	"UW"	estd	1934
------	------	------	------

e1

since	2015
-------	------

e11

since	1992
-------	------

e2

since	2006
-------	------

...

e4

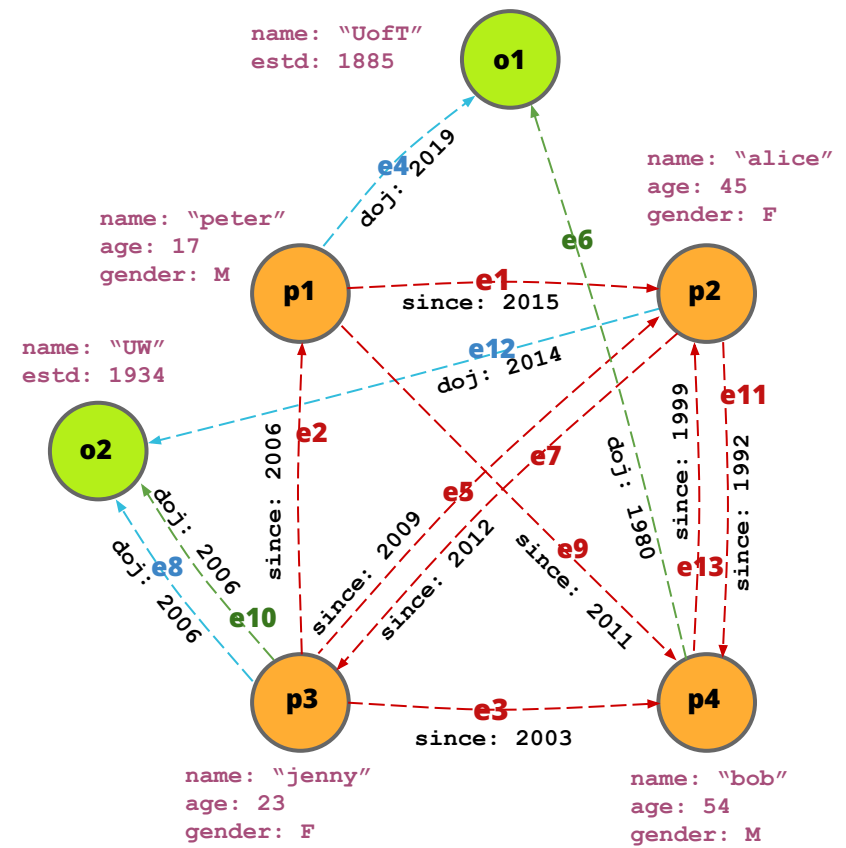
doj	2019
-----	------

e10

doj	2006
-----	------

...

Graph Databases 101: Query Processing



Cypher query:

```
MATCH
(a:PERSON) - [e:FOLLOWS] -> (b:PERSON)
WHERE
a.age > 30 & e.since > 2000
RETURN ...
```

subgraph patterns

constraints

Operators for query execution:

- Scan
- Filter
- Join
- GroupBy and Aggregate

Query plan:

[Scan a] → [Filter on a.age] → [Join a with b]
→ [Filter on e.since] → RETURN ...

Column-oriented RDBMS

Introduces a plethora of techniques to make analytical queries fast.

- Columnar storage

Positional offset based access to values in the column

p4	54	M
p2	23	F
p3	45	F
p1	17	F
	PERSON.age	PERSON.gender

- Columnar compression on homogenous data
- Late Materialization and Operations over Compressed Data
- Block-based Processing

Good cache locality. **Not optimized for many-to-many joins**

0111b

M=0, F=1

Dictionary-encoded PERSON gender

Motivation

Workloads on GDBMSs and Columnar RDBMS are similar but have fundamentally different access patterns.

Can we push the storage and performance limits of GDBMs by using the techniques that are tailored for achieving high performance in columnar RDBMSs ?

Sub questions:

- **What are the specific requirements for designing the storage layer of GDBMSs ?**
- **Can the existing columnar techniques directly apply to various components of the GDBMSs ?**
- **Can we present new techniques where existing techniques do not directly apply ?**
- **How can we adapt query processor to perform better with many-to-many joins ?**